

Lecture 8


Linear Regression (3)

Probabilistic view of linear regression

EECS 189/289, Fall 2025 @ UC Berkeley

Joseph E. Gonzalez and Narges Norouzi

Do not edit
How to change the
design



**Join at slido.com
#4234276**

① The Slido app must be installed on every computer you're presenting from

slido

Roadmap

- Recap of Regularization
- When Normal Equation Gets Tricky
- Least-Squares $\hat{=}$ Maximum Likelihood
- ~~Uncertainty $\approx$ Confidence Intervals~~
- ~~Decision Theory~~

Didn't get to cover in the lecture, not in scope for this lecture



4234276

Recap of Regularization

- Recap of Regularization
- When Normal Equation Gets Tricky
- Least-Squares $\hat{=}$ Maximum Likelihood



4234276

L2 Regularization (Ridge)



4234276

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{D}}[\mathbf{w}]} + \underbrace{\lambda \mathbf{E}_{\mathbf{w}}[\mathbf{w}]}$$

$$\frac{1}{2} \sum_{n=1}^N (\mathbf{t}_n - y(\mathbf{x}_n, \mathbf{w}))^2$$


$$\frac{1}{2} L_2[\mathbf{w}]^2 = \frac{1}{2} \sum_d \mathbf{w}_d^2$$

$$\frac{1}{2} (\|\mathbf{t}_n - y(\mathbf{x}_n, \mathbf{w})\|_2)^2$$

$$\frac{1}{2} (\|e_n\|_2)^2$$

length of the
residual vector

$$L_2[\mathbf{w}] = \sqrt{\sum_d \mathbf{w}_d^2}$$




L2 Regularization Demo

[Interactive link](#)



4234276





4234276

L1 Regularization (Lasso)

$$L_1[\mathbf{w}] = \sum_j |\mathbf{w}_j|$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{D}}[\mathbf{w}]} + \underbrace{\lambda \mathbf{E}_{\mathbf{w}}[\mathbf{w}]}$$


$$\frac{1}{2} \sum_{n=1}^N (\mathbf{t}_n - y(\mathbf{x}_n, \mathbf{w}))^2$$

$$\frac{1}{2} L_1[\mathbf{w}] = \frac{1}{2} \sum_d |\mathbf{w}_d|$$

$$\frac{1}{2} (\|\mathbf{t}_n - y(\mathbf{x}_n, \mathbf{w})\|_2)^2$$

$$\frac{1}{2} (\|e_n\|_2)^2$$

length of the
residual vector




L1 Regularization Demo

[Interactive link](#)



4234276



When Normal Equation Gets Tricky

- Recap of Regularization
- When Normal Equation Gets Tricky
- Least-Squares $\hat{=}$ Maximum Likelihood



4234276



4234276

Normal Equation (Reminder)

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{t}_n - y(\mathbf{x}_n, \mathbf{w}))^2 = \frac{1}{2} (\|\vec{e}\|_2)^2$$

$E(\mathbf{w})$ is minimum when $\vec{e}$ is orthogonal projection of $\mathbf{t}$.

Vector $\vec{e}$ is orthogonal to the $\text{span}(\mathbf{X})$, meaning that:

$$\mathbf{X}^T \vec{e} = 0$$

Adding the definition of residual

$$\mathbf{X}^T (\mathbf{t} - \mathbf{Y}(\mathbf{X}, \mathbf{w}^*)) = 0$$

Adding formula for $\mathbf{Y}(\mathbf{X}, \mathbf{w}^*) = \mathbf{X}\mathbf{w}^*$

$$\mathbf{X}^T \mathbf{t} - \mathbf{X}^T \mathbf{X} \mathbf{w}^* = 0$$


Moving terms

$$\mathbf{X}^T \mathbf{t} = \mathbf{X}^T \mathbf{X} \mathbf{w}^*$$

Normal
Equation

If $\mathbf{X}^T \mathbf{X}$ is
invertible

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$$





4234276

When Normal Equation Gets Tricky

- When is $\mathbb{X}^T \mathbb{X}$ invertible?

- Full column rank: $\text{rank}(\mathbb{X}) = D + 1$
- Requires $N \geq D + 1$ (at least as many samples as features)
- No perfect collinearity: $\det(\mathbb{X}^T \mathbb{X}) \neq 0$
- All singular values must be non-zero: $\sigma_1, \sigma_2, \dots, \sigma_{D+1} > 0$.

Normal Equation

$$\mathbf{w}^* = (\underbrace{\mathbb{X}^T \mathbb{X}}_{D+1})^{-1} \mathbb{X}^T \mathbf{t}$$
$$\mathbb{X} = \left[\begin{array}{c|ccc} 1 & \mathbf{x}_{11} & \cdots & \mathbf{x}_{1D} \\ 1 & \cdots & \cdots & \cdots \\ 1 & \mathbf{x}_{N1} & \cdots & \mathbf{x}_{ND} \end{array} \right] \left. \vphantom{\begin{array}{c} 1 \\ 1 \\ 1 \end{array}} \right\} N$$



If $N = D + 1$ and X is full rank,
what is the training error?

Solution



4234276

- If $N = D + 1$ and $\mathbb{X}$ is full rank, what is the training error?

$$\mathbb{X} \in \mathbb{R}^{N \times D+1}$$

$$\mathbb{X} \in \mathbb{R}^{D+1 \times D+1} \quad \text{Square matrix that is also full rank}$$

We can solve directly $\mathbb{X}^T \vec{e} = 0$

Multiply both sides by $(\mathbb{X}^T)^{-1}$ $\vec{e} = 0$

Training error will always be 0.



In which of the following situations is $X^T X$ likely to be singular or nearly singular?



4234276

When Normal Equation Gets Tricky

- When is $\mathbb{X}^T \mathbb{X}$ invertible?

- Full column rank: $\text{rank}(\mathbb{X}) = D + 1$
- Requires $N \geq D + 1$ (at least as many samples as features)
- No perfect collinearity: $\det(\mathbb{X}^T \mathbb{X}) \neq 0$
- All singular values must be non-zero: $\sigma_1, \sigma_2, \dots, \sigma_{D+1} > 0$.

Normal Equation

$$\mathbf{w}^* = (\underbrace{\mathbb{X}^T \mathbb{X}}_{D+1})^{-1} \mathbb{X}^T \mathbf{t}$$
$$\mathbb{X} = \left[\begin{array}{cccc} 1 & x_{11} & \cdots & x_{1D} \\ 1 & \cdots & \cdots & \cdots \\ 1 & x_{N1} & \cdots & x_{ND} \end{array} \right] \quad \left. \vphantom{\begin{array}{c} 1 \\ 1 \\ 1 \end{array}} \right\} N$$

- Ill-conditioning (numerical instability): Even if $\det(\mathbb{X}^T \mathbb{X}) \neq 0$

For any matrix A , $\det(A) = \prod_{i=1}^K \sigma_i$ where σ_i are eigenvalues.

When $\det(A) \approx 0$, means at least one singular value $\sigma_i \approx 0$.





4234276

When Normal Equation Gets Tricky

- When is $\mathbb{X}^T \mathbb{X}$ invertible?

- Full column rank: $\text{rank}(\mathbb{X}) = D + 1$
- Requires $N \geq D + 1$ (at least as many samples as features)
- No perfect collinearity: $\det(\mathbb{X}^T \mathbb{X}) \neq 0$
- All singular values must be non-zero: $\sigma_1, \sigma_2, \dots, \sigma_{D+1} > 0$.

Normal Equation

$$\mathbf{w}^* = (\underbrace{\mathbb{X}^T \mathbb{X}}_{D+1})^{-1} \mathbb{X}^T \mathbf{t}$$
$$\mathbb{X} = \left[\begin{array}{c|ccc} 1 & x_{11} & \cdots & x_{1D} \\ 1 & \cdots & \cdots & \cdots \\ 1 & x_{N1} & \cdots & x_{ND} \end{array} \right] \left. \vphantom{\begin{array}{c} 1 \\ 1 \\ 1 \end{array}} \right\} N$$

- Ill-conditioning (numerical instability): Even if $\det(\mathbb{X}^T \mathbb{X}) \neq 0$

- **Condition number** $\kappa = \frac{\sigma_{\max}}{\sigma_{\min}} \gg 10^8$ small noise in the data creates huge weight swing
- Why?
 - Rounding error will be amplified when calculating $\kappa \varepsilon_{\text{machine}}$.



4234276

Fix #1: Ridge Trick

- In Ridge regression, the closed form solution is

$$w^* = (X^T X + \lambda I)^{-1} X^T t$$

Proof in
discussion 4
this week

- Ridge replaces $X^T X$ in normal equation with $X^T X + \lambda I$.



4234276

Fix #1: Ridge Trick

- In Ridge regression, the closed form solution is

$$w^* = (X^T X + \lambda I)^{-1} X^T t$$

Proof in
discussion 4
this week

- Ridge replaces $X^T X$ in normal equation with $X^T X + \lambda I$.

What are the eigenvalues of $X^T X + \lambda I$?

- A. σ_i
- B. $\sigma_i + \lambda$
- C. σ_i^2
- D. $\sigma_i^2 + \lambda$



**Which of the following options
are correct?**



4234276

Fix #1: Ridge Trick

- In Ridge regression, the closed form solution is

$$w^* = (X^T X + \lambda I)^{-1} X^T t$$

Proof in
discussion 4
this week

- Ridge replaces $X^T X$ in normal equation with $X^T X + \lambda I$.
- This adds λ to every eigenvalue of $X^T X$:

$$\sigma_i^2 \rightarrow \sigma_i^2 + \lambda$$

- $\lambda > 0$ boosts all eigenvalues.
- Condition number shrinks:

$$\kappa_{ridge} = \frac{\sigma_{max}^2 + \lambda}{\sigma_{min}^2 + \lambda} \ll \kappa$$

This stabilizes
inversion

Illconditioning Example

With the condition factor that is large, due to very small eigenvalues resulting from very similar features, the weights become unstable.

Try changing the noise and you will see substantial swings in weights.

```
N, D = 500, 2 # one bias + 2 features
X = np.ones((N, D + 1)) # first column = bias
x1 = np.random.randn(N)
x2 = x1 + 1e-4 * np.random.randn(N) # almost identical feature
X[:, 1] = x1
X[:, 2] = x2
```

Condition $(X^T X) =$
420140223.27925146

```
t = 3 + 2*x1 - 1*x2 + 0.1*np.random.randn(N) # targets
```

```
XtX = X.T @ X
Xty = X.T @ t
w_ne = np.linalg.solve(XtX, Xty)
```

```
lam = 1e-2
w_ridge = np.linalg.solve(XtX + lam*np.eye(D+1), Xty)
```

```
w_neq: [ 3.00273423 12.87110771 -11.87562795]
```

```
w_ridge : [3.00275573 0.50070697 0.49481053]
```



4234276

How to resolve the issue of
ill-conditioned design matrix?



4234276

SVD Recap: Stretch + Rotate


- Every real matrix $\mathbb{X}$ can be factored as

$$\mathbb{X} = U\Sigma V^T$$

U and V are rotations and have orthogonal columns ($U^T U = V^T V = I$).

Σ is diagonal scaling by singular values σ_i .

$\Sigma = \text{diag}(\sigma_1, \dots, \sigma_r)$ with $r = \text{rank}(\mathbb{X})$.





4234276

Fix #2: Flipping Singular Values with Moore–Penrose Pseudo-Inverse

- Ordinary inverse of $\mathbb{X} = U\Sigma V^T$ is $\mathbb{X}^{-1} = V\Sigma^{-1}U^T$

- If some $\sigma_i = 0$, we cannot invert directly.

- Moore–Penrose Pseudo-inverse:

$$\sigma_{ii}^+ = \begin{cases} \frac{1}{\sigma_i} & \text{if } \sigma_i > \varepsilon \sigma_{max} \\ 0 & \text{otherwise} \end{cases}$$

- We safely “flip” only the non-tiny singular values.

- Pseudo-Inverse:

$$\mathbb{X}^+ = V\Sigma^+U^T$$



4234276

Fix #2: Flipping Singular Values with Moore–Penrose Pseudo-Inverse

- Pseudo-Inverse:

$$\mathbb{X}^+ = V\Sigma^+U^T$$

- Properties:

- $\mathbb{X}\mathbb{X}^+\mathbb{X} = \mathbb{X}$,
- $\mathbb{X}^+\mathbb{X}\mathbb{X}^+ = \mathbb{X}^+$,
- $(\mathbb{X}\mathbb{X}^+)^T = \mathbb{X}\mathbb{X}^+$,
- $(\mathbb{X}^+\mathbb{X})^T = \mathbb{X}^+\mathbb{X}$.

- These guarantee uniqueness of $\mathbb{X}^+$.



4234276

$\mathbb{X}\mathbb{X}^+$ Is the Orthogonal Projector

• Let $P = \mathbb{X}\mathbb{X}^+$. Then:

1. **Projector:** $P^2 = \mathbb{X}\mathbb{X}^+\mathbb{X}\mathbb{X}^+ = \mathbb{X}\mathbb{X}^+ = P$
2. **Symmetric:** $P^T = P$
3. **Range:** $Py = \mathbb{X}\mathbb{X}^+\mathbb{X}w = \mathbb{X}w = y$

P acts as identity on $\text{span}(\mathbb{X})$.
 $\text{range}(P)$ is $\text{span}(\mathbb{X})$.

$P = \mathbb{X}\mathbb{X}^+$ is the orthogonal projector onto $\text{span}(\mathbb{X})$.



4234276

Obtaining w^*

$P = \mathbb{X}\mathbb{X}^+$ is the orthogonal projector onto $\text{span}(\mathbb{X})$.



We will use P to find the projection of t onto the $\text{span}(\mathbb{X})$.

$$y = Pt$$

Replacing $P = \mathbb{X}\mathbb{X}^+$

$$y = \mathbb{X}\mathbb{X}^+t$$

Knowing that $y = \mathbb{X}w$

$$w^* = \mathbb{X}^+t$$

Very important result!!!

This derivation
is incorrect, see
next slide



4234276

Obtaining w^*

$P = \mathbb{X}\mathbb{X}^+$ is the orthogonal projector onto $\text{span}(\mathbb{X})$.



We will use P to find the projection of t onto the $\text{span}(\mathbb{X})$.

$$y = Pt \quad \text{Replacing } P = \mathbb{X}\mathbb{X}^+$$

$$y = \mathbb{X}\mathbb{X}^+t \quad \text{Knowing that } y = \mathbb{X}w$$

$$\mathbb{X}w = \mathbb{X}\mathbb{X}^+t \longrightarrow w = \mathbb{X}^+t + (I - \mathbb{X}\mathbb{X}^+)z \text{ for } z \in R^{D+1}$$

Using
minimum-norm
argument

$$w^* = \mathbb{X}^+t \quad \text{Very important result!!!}$$

$$w^* = \arg \min_{w: \mathbb{X}w=y} \|w\|_2^2$$



4234276

Fix #2: Moore–Penrose Pseudo-Inverse

How can Penrose Pseudo-Inverse help us with solving normal equation?

- Moore–Penrose Pseudo-inverse:

$$\sigma_{ii}^+ = \begin{cases} \frac{1}{\sigma_i} & \text{if } \sigma_i > \varepsilon \sigma_{max} \\ 0 & \text{otherwise} \end{cases}$$

- We safely “flip” only the non-tiny singular values.

Closed-form solution for least squares:

X .

$$\mathbf{w}^* = \mathbf{X}^+ \mathbf{t}$$

This works no matter whether $\mathbf{X}$ is tall ($N > D$), square, or wide ($N < D$). Also regardless of singular values.



4234276

Obtaining w^*

$P = \mathbb{X}\mathbb{X}^+$ is the orthogonal projector onto $\text{span}(\mathbb{X})$.



We will use P to find the projection of t onto the $\text{span}(\mathbb{X})$.

$$y = Pt \quad \text{Replacing } P = \mathbb{X}\mathbb{X}^+$$

$$y = \mathbb{X}\mathbb{X}^+t \quad \text{Knowing that } y = \mathbb{X}w$$

$$\mathbb{X}w = \mathbb{X}\mathbb{X}^+t \longrightarrow w = \mathbb{X}^+t + (I - \mathbb{X}\mathbb{X}^+)z \text{ for } z \in R^{D+1}$$

Using
minimum-norm
argument

$$w^* = \mathbb{X}^+t \quad \text{Very important result!!!}$$

$$w^* = \arg \min_{w: \mathbb{X}w=y} \|w\|_2^2$$



4234276

Why This w^* Amongst So Many?

If X is rank-deficient, there are infinitely many w with the same prediction y .

You can add any vector from the nullspace of X .

The Pseudo-inverse picks the one with smallest Euclidean form:

$$w^* = \arg \min_{w: Xw=y} \|w\|_2^2$$

Fix 2: Why It Helps?



4234276

Issue	How pseudo-inverse fixes it?
Rank deficiency (perfect multicollinearity)	



4234276

Fix 2: Why It Helps?

Issue	How pseudo-inverse fixes it?
Rank deficiency (perfect multicollinearity)	
Ill-conditioning	Truncation threshold discards numerically meaningless directions, implicitly adds ridge-like regularization.



4234276

Fix 2: Why It Helps?

Issue	How pseudo-inverse fixes it?
Rank deficiency (perfect multicollinearity)	
Ill-conditioning	Truncation threshold discards numerically meaningless directions, implicitly adds ridge-like regularization.



4234276

Fix 2: Why It Helps?

Issue	How pseudo-inverse fixes it?
Rank deficiency (perfect multicollinearity)	
Ill-conditioning	Truncation threshold discards numerically meaningless directions, implicitly adds ridge-like regularization.
Implementation convenience	

Least-Squares $\hat{=}$ Maximum Likelihood

- Recap of Regularization
- When Normal Equation Gets Tricky
- Least-Squares $\hat{=}$ Maximum Likelihood
- Uncertainty $\approx$ Confidence Intervals
- Decision Theory



4234276



4234276

Least Squares $\hat{=}$ Maximum Likelihood

Recall: We fit a linear featurized function by *Minimizing the sum-of-squares error*.

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{t}_n - y(\mathbf{x}_n, \mathbf{w}))^2$$

Is this just algebraic convenience ... or is there a probabilistic justification?



4234276

Additive-Gaussian Noise Assumption


- Assume the target t is the true signal plus noise:

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon$$

where $\epsilon \sim \mathcal{N}(0, \sigma^2)$ (zero-mean Gaussian random variable) $\rightarrow$ i.i.d. across points.

- The conditional likelihood of one observation is:

$$p(t | \mathbf{x}, \mathbf{w}, \sigma^2) = \mathcal{N}(t | y(\mathbf{x}, \mathbf{w}), \sigma^2)$$





4234276

Maximum Likelihood: Over the Dataset

- Consider the data set of i.i.d. inputs $\mathbb{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ with corresponding target values $t_1, \dots, t_N$. We group the target variables $\{t_n\}$ into a column vector that we denote by $\mathbf{t}$.

$$p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2) = \mathcal{N}(\mathbf{t}|y(\mathbb{X}, \mathbf{w}), \sigma^2)$$

Datapoints are
independent

$$= \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \sigma^2)$$

$$p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \sigma^2)$$



4234276

Least Squares $\hat{=}$ Maximum Likelihood

$$p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \sigma^2)$$

The optimum value of $\mathbf{w}$ in the above equation equals the optimum value of the log of the function (because log is a monotonically increasing function)

$$\ln(p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2)) = \sum_{n=1}^N \ln(\mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \sigma^2))$$

Writing the equation for a Normal distribution

$$= \sum_{n=1}^N \ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2} \right)$$

Applying $\log a \times b = \log a + \log b$

$$= \sum_{n=1}^N \left(\ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) - \frac{1}{2\sigma^2} (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2 \right)$$

Simplify and separate two terms


$$\ln(p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2)) = -\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2$$



4234276

Least Squares $\hat{=}$ Maximum Likelihood

$$\begin{aligned}\ln(p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2)) &= -\frac{N}{2}\ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2 \\ &= -\frac{N}{2} \ln(2\pi) - \frac{N}{2} \ln(\sigma^2) - \frac{1}{\sigma^2} E_D(\mathbf{w}) \quad (*)\end{aligned}$$


$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2$$

- The first two terms in can be treated as constants when determining $\mathbf{w}$ because they are independent of $\mathbf{w}$.
- Therefore, maximizing the likelihood function under a Gaussian noise distribution is equivalent to minimizing the sum-of-squares error function.

Maximum Likelihood



4234276

$$\nabla_{\mathbf{w}} \ln(p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2)) = -\frac{1}{\sigma^2} \nabla_{\mathbf{w}} E_{\mathbf{D}}(\mathbf{w})$$

Sum rule

$$= \nabla_{\mathbf{w}} \left(-\frac{1}{2\sigma^2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2 \right)$$

Chain rule

$$= -\frac{1}{2\sigma^2} \left(\sum_{n=1}^N \nabla_{\mathbf{w}} (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2 \right)$$

Simplify

$$\begin{aligned} &= -\frac{1}{2\sigma^2} 2 \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n)) (-\phi(\mathbf{x}_n)) \\ &= \frac{1}{\sigma^2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n) \end{aligned}$$

Setting the gradient to 0 gives:

$$\frac{1}{\sigma^2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n) = 0$$

$$\sum_{n=1}^N t_n \phi(\mathbf{x}_n) - \mathbf{w}^T \phi(\mathbf{x}_n) \phi(\mathbf{x}_n) = 0$$

$$\sum_{n=1}^N t_n \phi(\mathbf{x}_n) = \sum_{n=1}^N \boxed{\mathbf{w}^T \phi(\mathbf{x}_n)} \phi(\mathbf{x}_n) \quad \text{const}$$

$$\sum_{n=1}^N t_n \phi(\mathbf{x}_n) = \sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T \mathbf{w}$$

Maximum Likelihood



4234276

$$\sum_{n=1}^N t_n \phi(\mathbf{x}_n) = \sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T \mathbf{w}$$

If we stack features $\phi(\mathbf{x}_n)$ row-wise in the design matrix $\Phi(\mathbb{X})$

$$\Phi(\mathbb{X}) = \begin{bmatrix} \phi(\mathbf{x}_1)^T \\ \phi(\mathbf{x}_2)^T \\ \dots \\ \phi(\mathbf{x}_N)^T \end{bmatrix}$$

$$\Phi(\mathbb{X})^T \mathbf{t} = \Phi(\mathbb{X})^T \Phi(\mathbb{X}) \mathbf{w}_{ML}$$

$$\mathbf{w}_{ML} = (\Phi(\mathbb{X})^T \Phi(\mathbb{X}))^{-1} \Phi(\mathbb{X})^T \mathbf{t}$$

Normal Equations for the least squared problem



4234276

Least Squares $\hat{=}$ Maximum Likelihood

$$p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \sigma^2)$$

$$\ln(p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2)) = -\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} E_D(\mathbf{w})$$

$$\text{with } E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \phi(\mathbf{x}_n))^2$$

$$\nabla_{\mathbf{w}} \ln(p(\mathbf{t}|\mathbb{X}, \mathbf{w}, \sigma^2)) = -\frac{1}{\sigma^2} \nabla_{\mathbf{w}} E_D(\mathbf{w}) \longrightarrow \mathbf{w}_{ML} = (\Phi(\mathbb{X})^T \Phi(\mathbb{X}))^{-1} \Phi(\mathbb{X})^T \mathbf{t}$$

$$\Phi(\mathbb{X}) = \begin{bmatrix} \phi(\mathbf{x}_1)^T \\ \phi(\mathbf{x}_2)^T \\ \vdots \\ \phi(\mathbf{x}_N)^T \end{bmatrix}$$

Hence the least-squares solution is the **MLE** under Gaussian noise.

Lecture 8

Linear Regression (3)

Credit: Joseph E. Gonzalez and Narges Norouzi

Reference Book Chapters: Chapter sections 4.1 and 4.2