

Conv. + Momentum + Adam + Stochastic Gradient Descent

The core algorithm of modern machine learning

EECS 189/289, Fall 2025 @ UC Berkeley

Joseph E. Gonzalez and Narges Norouzi

Do not edit
How to change the
design



**Join at slido.com
#3389863**

① The Slido app must be installed on every computer you're presenting from

slido



The (Batch) Gradient Descent Alg.

Update the weights by moving in the **opposite direction** of the gradient:

```
 $\mathbf{w}^{(0)} = \text{Initialize}()$   
for  $\tau$  in range(1,  $\tau_{\text{max}}$ ):  
     $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} - \eta \nabla E(\mathbf{w}^{(\tau-1)})$   
    if  $\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ : terminate
```


$\text{Initialize}()$ typically a small random value (more on this later)

$\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ stopping condition (e.g., ϵ change in $\Delta \mathbf{w}^{(\tau-1)}$)

Hyper Parameters:

- η is the learning rate (typically a small value $0 < \eta < 1$, more on this later)
- τ_{max} is the maximum number of iterations

What are the conditions for convergence?



—●— Gradient Descent Path

Learning rate
seemed to affect
convergence even on
a simple parabola.

Convergence Assuming Quadratic Approximation



Quadratic Approximations

The 2nd order (quadratic) Taylor approximation of a function $f: \mathbb{R} \rightarrow \mathbb{R}$ evaluated at $\hat{x}$ is:

$$\tilde{f}(x) = f(\hat{x}) + \left(\frac{\partial}{\partial x} f(x) \Big|_{x=\hat{x}} \right) (x - \hat{x}) + \frac{1}{2} \left(\frac{\partial^2}{\partial x^2} f(x) \Big|_{x=\hat{x}} \right) (x - \hat{x})^2$$

For example: $f(x) = \sin(x)$

$$\begin{aligned} \tilde{f}_{\hat{x}=1}(x) &= \sin(1) + \cos(1)(x - 1) \\ &\quad + \frac{1}{2}(-\sin(1))(x - 1)^2 \\ \tilde{f}_{\hat{x}=5}(x) &= \sin(5) + \cos(5)(x - 5) \\ &\quad + \frac{1}{2}(-\sin(5))(x - 5)^2 \end{aligned}$$



Taylor Expansion of the Error Surface

The 2nd-order Taylor expansion of the **error surface** at $\hat{w}$:

$$\tilde{E}_{\hat{w}}(w) = E(\hat{w}) + (w - \hat{w})^T b + \frac{1}{2} (w - \hat{w})^T H (w - \hat{w})$$

- $b = \nabla E(\hat{w})$ is the **Gradient** of E evaluated at $\hat{w}$ (vector)
- $H = \nabla \nabla E(\hat{w})$ is the **Hessian** of E evaluated at $\hat{w}$ (a matrix)

Matrix Equivalent of
the Second Derivative



Hessians (the 2nd derivative)

The Hessian of a **scalar-valued function** $E: \mathbb{R}^D \rightarrow \mathbb{R}$ is the **matrix-valued function** $\nabla^2 E: \mathbb{R}^D \rightarrow \mathbb{R}^{D \times D}$ of the 2nd partial derivatives of f

$$H = \nabla^2 E = \begin{bmatrix} \frac{\partial^2}{\partial x_1^2} E(x) & \dots & \frac{\partial^2}{\partial x_1 x_D} E(x) \\ \vdots & \ddots & \vdots \\ \frac{\partial^2}{\partial x_D x_1} E(x) & \dots & \frac{\partial^2}{\partial x_D^2} E(x) \end{bmatrix}$$

The Hessian

- describes the local **curvature** of E
- determines the **critical point type** (minima, maxima, or saddle)



Hessian Exercise

Consider the function:

$$E(w) = (w_0 - 1)^2 + (w_1 - 2)^2 + 1$$

Then the gradient is: $\nabla E(w) = \begin{bmatrix} 2(w_0 - 1) \\ 2(w_1 - 2) \end{bmatrix}$

The Hessian is:

$$H(w) = \begin{bmatrix} \frac{\partial^2}{\partial w_0^2} E(w) & \frac{\partial^2}{\partial w_0 \partial w_1} E(w) \\ \frac{\partial^2}{\partial w_0 \partial w_1} E(w) & \frac{\partial^2}{\partial w_1^2} E(w) \end{bmatrix} = \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix}$$

Hessian

$$\nabla^2 E = \begin{bmatrix} \frac{\partial^2}{\partial x_1^2} E(x) & \dots & \frac{\partial^2}{\partial x_1 \partial x_D} E(x) \\ \vdots & \ddots & \vdots \\ \frac{\partial^2}{\partial x_D \partial x_1} E(x) & \dots & \frac{\partial^2}{\partial x_D^2} E(x) \end{bmatrix}$$



Taylor Expansion of the Error Surface

The 2nd-order Taylor expansion of the error surface at $\hat{w}$:

$$\tilde{E}_{\hat{w}}(w) = E(\hat{w}) + (w - \hat{w})^T b + \frac{1}{2} (w - \hat{w})^T H (w - \hat{w})$$

- $b = \nabla E(\hat{w})$ is the **Gradient** of E evaluated at $\hat{w}$ (vector)
- $H = \nabla \nabla E(\hat{w})$ is the **Hessian** of E evaluated at $\hat{w}$ (a matrix)

If $\hat{w}$ is a **critical point** (w^*) then $\nabla E(w^*) = 0 = b$

$$\tilde{E}_{w^*}(w) = E(w^*) + \frac{1}{2} (w - w^*)^T H (w - w^*)$$

Use a change of variables to better understand this part.



Eigen Decomposition of the Hessian

We can construct the **Eigen decomposition of the Hessian**

$$Hu_i = \lambda_i u_i$$

- u_i are the complete set of **orthonormal eigenvectors**: $\forall_{i,j}: u_i^T u_j = \delta_{ij}$
- λ_i are the real-valued **eigenvalues**

Then we apply a **change of variables**:

$$w - w^* = \sum_i \alpha_i u_i \quad \text{where} \quad \alpha_i = (w - w^*)^T u_i$$

This allows us to rewrite:

$$\tilde{E}_{w^*}(w) = E(w^*) + \frac{1}{2} (w - w^*)^T H (w - w^*) = E(w^*) + \frac{1}{2} \sum_i \alpha_i^2 \lambda_i$$

Show it!



- u_i are the complete set of **orthonormal eigenvectors**: $\forall_{i,j}: u_i^T u_j = \delta_{ij}$
- λ_i are the real-valued **eigenvalues**

Then we apply a **change of variables**:

$$w - w^* = \sum_i \alpha_i u_i \quad \text{where} \quad \alpha_i = (w - w^*)^T u_i$$

This allows us to rewrite:

$$\tilde{E}_{w^*}(w) = E(w^*) + \frac{1}{2} (w - w^*)^T H (w - w^*) = E(w^*) + \frac{1}{2} \sum_i \alpha_i^2 \lambda_i$$

Show it! 

Derivation:

$$\begin{aligned} (w - w^*)^T H (w - w^*) &= \left(\sum_i \alpha_i u_i \right)^T H \left(\sum_j \alpha_j u_j \right) = \left(\sum_i \alpha_i u_i \right)^T \left(\sum_j \alpha_j H u_j \right) \\ &= \left(\sum_i \alpha_i u_i \right)^T \left(\sum_j \alpha_j \lambda_j u_j \right) = \sum_{ij} \alpha_i \alpha_j u_i^T u_j \lambda_j = \sum_{ij} \alpha_i \alpha_j \delta_{ij} \lambda_j = \sum_i \alpha_i^2 \lambda_i \end{aligned}$$

Eigenvector Defn.

Orthonormality



Taylor Expansion at the Stationary Pts.

If $\hat{w}$ is a **stationary point** (w^*) (where $b = \nabla E|_{w=\hat{w}} = 0$)

$$\tilde{E}_{w^*}(\alpha) = E(w^*) + \frac{1}{2} \sum_i \alpha_i^2 \lambda_i$$

At $\alpha = 0$ we get the error at the critical point $E(w^*)$.

- If we move in **direction i** (corresponding to u_i) then α_i^2 **increases**.
 - The change in error then depends on the sign of λ_i

Therefore, if the **eigenvalues** are

- **Slido!**



If the eigen values at w^* are all positive then



Taylor Expansion at the Stationary Pts.

If $\hat{w}$ is a **stationary point** (w^*) (where $b = \nabla E|_{w=\hat{w}} = 0$)

$$\tilde{E}_{w^*}(\alpha) = E(w^*) + \frac{1}{2} \sum_i \alpha_i^2 \lambda_i$$

At $\alpha = 0$ we get the error at the critical point $E(w^*)$.

- If we move in **direction i** (corresponding to u_i) then α_i^2 **increases**.
 - The change in error then depends on the sign of λ_i

Therefore, if the **eigenvalues** are

- **all positive** then w^* is a **local minimum**.
- **all negative** then w^* is a **local maximum**.
- **mixed** then w^* is a **saddle point**.



Stationary Points and the Hessian

Eigenvalues of the Hessian determine the curvature at the critical points

Minimum

Maximum

Saddle Point





**The eigenvector u_1
corresponds to the:**

Eigenvectors of the Hessian



Elliptical contours of constant error align with the eigenvectors of the Hessian matrix.





Demo

Taylor Expansions and the Hessian



3389863





Recap: Quadratic Approx. Error

We constructed the **2nd-order Taylor expansion** of the error surface at the **stationary point** w^* (where $b = \nabla E|_{w=w^*} = 0$):


$$\tilde{E}_{w^*}(w) = E(w^*) + \frac{1}{2} (w - w^*)^T H (w - w^*)$$

Using the **Eigen-decomposition of the Hessian**: $Hu_i = \lambda_i u_i$ we applied a change of variables $\alpha_i = (w - w^*)^T u_i$ to obtain:

$$\tilde{E}_{w^*}(\alpha) = E(w^*) + \frac{1}{2} \sum_i \alpha_i^2 \lambda_i$$

- This is minimized at $\alpha = 0$

We can **evaluate gradient descent** on this simple quadratic approximation to better understand the convergence properties.





Grad. Descent on the Quadratic Approx.

Compute the **gradient** with respect to α_i :

$$\frac{\partial}{\partial \alpha_i} \tilde{E}_{w^*}(\alpha) = \alpha_i \lambda_i$$

Quadratic Approx. Error

$$\tilde{E}_{w^*}(\alpha) = E(w^*) + \frac{1}{2} \sum_i \alpha_i^2 \lambda_i$$

Which gives the **gradient descent update**:

$$\alpha_i^{(\tau)} = \alpha_i^{(\tau-1)} - \eta \left(\alpha_i^{(\tau-1)} \lambda_i \right) = (1 - \eta \lambda_i) \alpha_i^{(\tau-1)}$$

If we **iteratively** apply the gradient descent update, we obtain:

$$\alpha_i^{(1)} = (1 - \eta \lambda_i) \alpha_i^{(0)} \rightarrow \alpha_i^{(2)} = (1 - \eta \lambda_i) \left((1 - \eta \lambda_i) \alpha_i^{(0)} \right) \rightarrow \dots \rightarrow$$

$$\alpha_i^{(\tau)} = (1 - \eta \lambda_i)^\tau \alpha_i^{(0)}$$



Grad. Descent on the Quadratic Approx.

Assuming a **quadratic approximation** to the error function around the minima, the τ^{th} iteration of gradient descent we have:

$$\alpha_i^{(\tau)} = (1 - \eta \lambda_i)^\tau \alpha_i^{(0)}$$

and the **error is minimized when $\alpha_i = 0$** .

To ensure $\alpha_i^{(\tau)} \rightarrow 0$ as $\tau \rightarrow \infty$ the learning rate should satisfy:

$$|1 - \eta \lambda_i| < 1 \quad (\text{Convergence Condition})$$

- Smaller $|1 - \eta \lambda_i|$ is better (converges faster)
- Negative $1 - \eta \lambda_i < 0$ oscillatory behavior



Learning Rate Impact on Convergence

To maximize the **progress on each step** of gradient descent we want to **make η as large as possible (take bigger steps)**

- We want $(1 - \eta\lambda_i)^\tau \rightarrow 0$ quickly as τ increases (larger $\eta \rightarrow$ smaller $1 - \eta\lambda_i$).
- However, to **ensure convergence we require $|1 - \eta\lambda_i| < 1$ for all λ_i**

Let $\lambda_{\min}$ and $\lambda_{\max}$ be the smallest and largest **eigenvalues** of the Hessian

- Then $\eta < \frac{2}{\lambda_{\max}}$ because if $\eta \geq \frac{2}{\lambda_{\max}}$ then $|1 - \eta\lambda_{\max}| \geq 1$

The **slowest converging dimension correspond to $\lambda_{\min}$** so the smallest we can make $(1 - \eta\lambda_i)$:

$$\left(1 - \frac{2}{\lambda_{\max}}\lambda_{\min}\right) < (1 - \eta\lambda_i)$$



Condition Number and Loss Surface


The slowest converging dimension (i) is the dimension corresponding to λ_{min} and its convergence to 0 is lower bounded by:

$$\left(1 - \frac{2}{\lambda_{max}} \lambda_{min}\right)^{\tau} \alpha_i^{(0)}$$

The ratio of $\frac{\lambda_{min}}{\lambda_{max}}$ is the **reciprocal** of the **condition number**.

- small $\frac{\lambda_{min}}{\lambda_{max}}$ implies slow convergence

Gradient descent converges quickly on **well conditioned problems**.





Issues with Gradient Descent (GD)

GD converges slowly when the **error surface is relatively flat**:

Step size
(learning rate η) is
too small.

GD oscillates when the error surface is **poorly conditioned**

Step size
(learning rate η) is
too big.



3389863

Demo

Oscillating convergence



Momentum

Momentum can be used to improve convergence in these settings by **accumulating common update directions**

$$\Delta \mathbf{w}^{(\tau-1)} = -\eta \nabla E(\mathbf{w}^{(\tau-1)}) + \mu \Delta \mathbf{w}^{(\tau-2)}$$

$$\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} + \Delta \mathbf{w}^{(\tau-1)}$$

- $0 \leq \mu < 1$ ($\mu = 0.9$) is the **momentum hyperparameter**

We now need to maintain the previous $\Delta \mathbf{w}^{(\tau-2)}$ update direction

Intuition: The solution “**picks up speed**” along the direction of repeated updates.



How Momentum Helps with Flat Dir.

In **directions** that **are relatively flat** (e.g., is approx. constant ∇E):

$$\begin{aligned}\Delta \mathbf{w}^{(\infty)} &= -\eta \nabla E + \mu \left(-\eta \nabla E + \mu \left(-\eta \nabla E + \mu (\dots) \right) \right) \\ \Delta \mathbf{w}^{(\infty)} &= -\eta \nabla E \underbrace{(1 + \mu + \mu^2 + \mu^3 + \dots)}_{\text{Geometric Series}} = -\left(\frac{\eta}{1 - \mu} \right) \nabla E\end{aligned}$$

The **momentum parameter** μ **increases** the effective learning rate

- Remember that $0 \leq \mu < 1$ and therefore $\eta \leq \left(\frac{\eta}{1 - \mu} \right)$



How Momentum Helps with Oscillations

In **directions** that **are steep** (e.g., oscillating sign ∇E):

$$\Delta \mathbf{w}^{(\tau)} = -\eta \nabla E + \mu \left(-\eta (-\nabla E) + \mu (-\eta \nabla E + \mu (\dots)) \right)$$




$$\Delta \mathbf{w}^{(\infty)} = -\eta \nabla E \underbrace{(1 - \mu + \mu^2 - \mu^3 + \dots)}_{\text{Alternating Geometric Series}} = -\left(\frac{\eta}{1 + \mu} \right) \nabla E$$

The **momentum parameter** μ **decrease** the effective learning rate

- Remember that $0 \leq \mu < 1$ and therefore $\left(\frac{\eta}{1 + \mu} \right) \leq \eta$

Demo

Momentum





Learning Rate Schedules

A larger initial learning rate can accelerate convergence but to ensure convergence we often reduce the learning rate across iterations.

- **Linear Schedule:** interpolates between η_{start} and η_{end} over T iterations

$$\eta^{(\tau)} = \left(1 - \frac{\tau}{T}\right) \eta_{\text{start}} + \left(\frac{\tau}{T}\right) \eta_{\text{end}}$$

- **Power-law Schedule:** (most common) with η_{start} , s , and c hyperparameters

$$\eta^{(\tau)} = \eta_{\text{start}} \left(1 + \frac{\tau}{s}\right)^{-c}$$

- **Exponential Schedule:** with η_{start} , s , and c hyperparameters

$$\eta^{(\tau)} = \eta_{\text{start}} c^{\lfloor \frac{\tau}{s} \rfloor}$$

- **Cosine Schedule:** common for LLMs, with η_{min} , η_{max} , and T hyperparameters


$$\eta^{(\tau)} = \eta_{\text{min}} + \frac{1}{2} (\eta_{\text{max}} - \eta_{\text{min}}) \left(1 + \cos\left(\frac{\tau}{T} \pi\right)\right)$$

Demo

Learning Rate Schedules



3389863





Adaptive Learning Rates – AdaGrad

AdaGrad* (Adaptive Gradient) reduces learning rate on dimensions with greatest curvature (largest gradient magnitude).

Reduce learning rate by running squared average of the gradient.

$$r_i^{(\tau)} = r_i^{(\tau-1)} + \left(\frac{\partial}{\partial w_i} E(\mathbf{w}) \right)^2$$

Tracking Magnitude
(terms squared)

$$w_i^{(\tau)} = w_i^{(\tau-1)} - \frac{\eta}{\sqrt{r_i^{(\tau)} + \epsilon}} \left(\frac{\partial}{\partial w_i} E(\mathbf{w}) \right)$$

Scaled Learning Rate

Issue:

Large initial gradients result in overly small learning rates near convergence.

- $\epsilon \approx 10^{-8}$ is a small constant for numerical stability.

*Developed at UC Berkeley



Adaptive Learning Rates – RMSProp

• **RMSProp** focuses on the **most recent gradients** by replacing the average of the gradients with an **exponential weighted average**

AdaGrad $r_i^{(\tau)} = r_i^{(\tau-1)} + \left(\frac{\partial}{\partial w_i} E(\mathbf{w}) \right)^2$

RMSProp $r_i^{(\tau)} = \beta r_i^{(\tau-1)} + (1 - \beta) \left(\frac{\partial}{\partial w_i} E(\mathbf{w}) \right)^2$

- Initialize $r_i^{(0)} = 0$.
- The **hyperparameter** β is typically $\beta = 0.9$



Adam = RMSProp + Momentum

Adam (adaptive moments) combines RMSProp with moments.

Momentum

$$s_i^{(\tau)} = \beta_1 s_i^{(\tau-1)} + (1 - \beta_1) \left(\frac{\delta}{\delta w_i} E(\mathbf{w}) \right)$$

RMSProp

$$r_i^{(\tau)} = \beta_2 r_i^{(\tau-1)} + (1 - \beta_2) \left(\frac{\delta}{\delta w_i} E(\mathbf{w}) \right)^2$$

Gradient Update

$$w_i^{(\tau)} = w_i^{(\tau-1)} - \eta \frac{s_i^{(\tau)}}{\sqrt{r_i^{(\tau)} + \delta}}$$

- Initialize $s_i^{(0)} = 0$ and $r_i^{(0)} = 0$.
- The hyperparameters β_1 and β_2 are typically $\beta_1 = 0.9$ and $\beta_2 = 0.99$.

Adam is probably the most widely used gradient descent update.

Bias Corrected Version:

$$\hat{s}_i^{(\tau)} = \frac{s_i^{(\tau)}}{1 - \beta_1^\tau}$$

$$\hat{r}_i^{(\tau)} = \frac{r_i^{(\tau)}}{1 - \beta_2^\tau}$$

$$w_i^{(\tau)} = w_i^{(\tau-1)} - \frac{\eta}{\sqrt{\hat{r}_i^{(\tau)} + \delta}} \hat{s}_i^{(\tau)}$$



Addresses issue where early steps are too small.

Demo

Adam



3389863





Recap: Batch Gradient Descent

Update the weights by moving in the **opposite direction** of the gradient:

```
 $\mathbf{w}^{(0)} = \text{Initialize}()$   
for  $\tau$  in range(1,  $\tau_{\text{max}}$ ):  
     $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} + \text{Step}(\nabla E(\mathbf{w}^{(\tau-1)}), \tau, \beta, \eta)$   
    if  $\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ : terminate
```

$\text{Initialize}()$ typically a small random value (more on this later)

$\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ stopping condition (e.g., ϵ change $\Delta \mathbf{w}^{(\tau-1)}$)

$\text{Step}(\nabla E(\mathbf{w}^{(\tau-1)}), \tau, \beta, \eta)$ defines how the gradient is used to compute $\Delta \mathbf{w}^{(\tau-1)}$

- This is where Adam would be implemented

Approximating Gradient Descent

Stochastic Gradient Descent



The Cost of Batch Gradient Descent

What is the cost of each iteration of batch gradient descent?

```
 $\mathbf{w}^{(0)} = \text{Initialize}()$   
for  $\tau$  in range(1,  $\tau_{\text{max}}$ ):  
     $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} + \text{Step}(\nabla E(\mathbf{w}^{(\tau-1)}), \tau, \beta, \eta)$   
    if  $\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ : terminate
```

If we have N training data points in D dimensions:

$$\nabla E(\mathbf{w}^{(\tau-1)}) = \nabla E(\mathbf{w}^{(\tau-1)}; \{(x_n, t_n)\}_{n=1}^N) = \frac{1}{N} \sum_{n=1}^N \nabla L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n))$$



What is the cost for a single gradient update assuming N data points in D dimensions and the basic gradient update equation.



The Cost of Batch Gradient Descent

What is the cost of each iteration of batch gradient descent?

```
 $\mathbf{w}^{(0)} = \text{Initialize}()$   
for  $\tau$  in range(1,  $\tau_{\max}$ ):  
     $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} + \text{Step}(\nabla E(\mathbf{w}^{(\tau-1)}), \tau, \beta, \eta)$   
    if  $\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ : terminate
```

If we have N training data points in D dimensions:

$$\nabla E(\mathbf{w}^{(\tau-1)}) = \nabla E(\mathbf{w}^{(\tau-1)}; \{(x_n, t_n)\}_{n=1}^N) = \frac{1}{N} \sum_{n=1}^N \nabla L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n))$$

The cost per update is $\mathcal{O}(ND)$

$\mathcal{O}(N)$

$\mathcal{O}(D)$



The Training Error Estimates the Test Error

As an **optimizer**, our goal has been to minimize the **average training error**:

$$E_{\text{train}}[\mathbf{w}] = \frac{1}{N} \sum_{n=1}^N L(t_n; f_{\mathbf{w}}(x_n))$$

However, we would have preferred to minimize the **expected test error**:

$$E_{\text{test}}[\mathbf{w}] = \mathbb{E}_{p(t_*, x_*)} [L(t_*, f_{\mathbf{w}}(x_*))]$$

- Where $p(t_*, x_*)$ is the **unknown distribution over test examples**.

If we assume that our training data is drawn from $p(t_*, x_*)$ then for large enough N :

$$E_{\text{train}}[\mathbf{w}] = \frac{1}{N} \sum_{n=1}^N L(t_n; f_{\mathbf{w}}(x_n)) \approx \mathbb{E}_{p(t_*, x_*)} [L(t_*, f_{\mathbf{w}}(x_*))] = E_{\text{test}}[\mathbf{w}]$$

The training error is an empirical estimate of the unknown test error.



Approximating the Gradient

The error (and the gradient) are **empirical expectations** of the **test error**:

$$E(\mathbf{w}^{(\tau-1)}) = \frac{1}{N} \sum_{n=1}^N L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n)) \approx \mathbb{E}_{p(t,x)} [L(t, f_{\mathbf{w}^{(\tau-1)}}(x))]$$

The error at a **single point** is *also* a (**noisy**) estimate:

$$E_n(\mathbf{w}^{(\tau-1)}) = L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n)) \approx \mathbb{E}_{p(t,x)} [L(t, f_{\mathbf{w}^{(\tau-1)}}(x))]$$

Compute the gradient for a **randomly (stochastically)** sampled **single point**:

$$\nabla E_n(\mathbf{w}^{(\tau-1)}) \approx \nabla E(\mathbf{w}^{(\tau-1)}; \{(x_n, t_n)\}_{n=1}^N) \approx \mathbb{E}_{p(t,x)} [\nabla L(t, f_{\mathbf{w}^{(\tau-1)}}(x))]$$

Cheap approx.

Best estimate (but expensive)

Ideal but don't know $p(t, x)$



Stochastic Gradient Descent

Gradient descent but with the gradient computed on a **single randomly sampled training data point**

```
w(0) = Initialize()  
for  $\tau$  in range(1,  $\tau_{\text{max}}$ ):  
    n = RandomInteger(0, N)  
    w( $\tau$ ) = w( $\tau-1$ ) + Step( $\nabla E_n$ (w( $\tau-1$ )),  $\tau$ ,  $\beta$ ,  $\eta$ )  
    if Converged(w( $\tau$ ), w( $\tau-1$ ),  $\epsilon$ ): terminate
```

Gradient of the error on the n^{th} data point.

$$\nabla E_n(\mathbf{w}^{(\tau-1)}) = \nabla L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n))$$

- Stochasticity comes from **random gradient estimates** $\nabla E_n(\mathbf{w}^{(\tau-1)})$
- **Cost per step is $\mathcal{O}(D)$** instead of $\mathcal{O}(ND)$ in batch gradient descent



Stochastic Gradient Descent (Shuffling)

In practice, we often **shuffle** and then **iterate** over the entire dataset

```
 $\mathbf{w}^{(0)} = \text{Initialize}();$ 
```

```
 $\tau = 1$ 
```

```
for  $e$  in range(0,  $e_{\text{max}}$ ): # Epochs
```

```
     $\text{ShuffleData}()$  # Randomize the Data Ordering
```

```
    for  $n$  in range(0,  $n$ ):
```

```
         $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} + \text{Step}(\nabla E_n(\mathbf{w}^{(\tau-1)}), \tau, \beta, \eta)$ 
```

```
        if  $\text{Converged}(\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon)$ : terminate
```

```
         $\tau = \tau + 1$ 
```

One **epoch** is a complete pass over all the data points.

Often only randomize data once.



Gradient of the error on the n^{th} data point.

Demo

Stochastic Gradient Descent




3389863





A Better Gradient Approximation

Estimating the gradient using a single data point results in **noisy estimates**.





A Spectrum of Stochastic Gradients

We typically refer to any of the sampling based gradient estimation procedures as **Stochastic Gradient Descent (SGD)**

Noisy but Cheap

Accurate but Expensive

$$\nabla L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n))$$

**Stochastic
Gradient Descent**
 $\mathcal{O}(D)$

$$\frac{1}{B} \sum_{b=1}^B \nabla L(t_{\mathcal{B}_b}, f_{\mathbf{w}^{(\tau-1)}}(x_{\mathcal{B}_b}))$$

**Mini-Batch Stochastic
Gradient Descent**
 $\mathcal{O}(BD)$

$$\frac{1}{N} \sum_{n=1}^N \nabla L(t_n, f_{\mathbf{w}^{(\tau-1)}}(x_n))$$

**Batch
Gradient Descent**
 $\mathcal{O}(ND)$

Stochastic Gradient Methods



Mini-batch Stochastic Gradient Descent

Gradient descent but with the gradient computed on a **batch of randomly sampled training points**

```
w(0) = Initialize()  
for  $\tau$  in range(1,  $\tau_{\text{max}}$ ):  
     $\mathcal{B}$  = RandBatch(0, N, size = B)  
     $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} + \text{Step}(\nabla E_{\mathcal{B}}(\mathbf{w}^{(\tau-1)}), \tau, \beta, \eta)$   
    if Converged( $\mathbf{w}^{(\tau)}, \mathbf{w}^{(\tau-1)}, \epsilon$ ): terminate
```

$$\begin{aligned} \nabla E_{\mathcal{B}}(\mathbf{w}^{(\tau-1)}) \\ &= \frac{1}{B} \sum_{b=1}^B \nabla L(t_{\mathcal{B}_b}, f_{\mathbf{w}^{(\tau-1)}}(x_{\mathcal{B}_b})) \end{aligned}$$

- **Cost per step is $\mathcal{O}(BD)$** instead of $\mathcal{O}(ND)$ in batch gradient descent
- Note that at batch size $B = 1$ **mini-batch** is the same as **SGD**
 - We often call both algorithms **SGD** (just with a **batch-size hyperparameter**)



Shuffling for Minibatch Gradient Descent

In practice, we often **shuffle** and then **iterate** over the entire dataset

```
 $w^{(0)} = \text{Initialize}();$ 
```

```
 $\tau = 1$ 
```

```
for  $e$  in range(0,  $e_{\max}$ ): # Epochs
```

```
     $\text{ShuffleData}()$  # Randomize the Data Ordering
```

```
    for  $n$  in range(0,  $n$ , step= $B$ ):
```

```
         $w^{(\tau)} = w^{(\tau-1)} + \text{Step}(\nabla E_{n, \dots, n+B}(w^{(\tau-1)}), \tau, \beta, \eta)$ 
```

```
        if  $\text{Converged}(w^{(\tau)}, w^{(\tau-1)}, \epsilon)$ : terminate
```

```
         $\tau = \tau + 1$ 
```

One **epoch** is a complete pass over all the data points.

Often only randomize data once.

Gradient of the error on the next B data points.



If I have N data points and I choose a batchsize $B=2$. In one epoch how many gradient steps will I take?



How do I Pick the Batch Size?

Larger batch size:

- **Better gradient estimate**
- **Enables hardware parallelism** – can compute gradient for each point in parallel

Smaller batch size:

- **Cheap to compute** □ take more steps
- **Greater stochasticity** can help with local minima

Typical design decision:



- Increase batch size to **saturate hardware parallelism**
- **Increase learning rate** proportional to batch size (take bigger steps)

Demo

Mini-batch Gradient Descent



3389863



Lecture 13

Stochastic Gradient Descent

Credit: Joseph E. Gonzalez and Narges Norouzi

Reference Book Chapters: Chapter 7