

Actor-Critic

CS 185/285

Instructor: Sergey Levine
UC Berkeley



Part 1:

Improving the policy gradient with values

Recap: policy gradients


REINFORCE algorithm:

- 
1. sample $\{\tau^{(i)}\}$ from $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$ (run the policy)
 2. $\nabla_\theta J(\theta) \approx \sum_i \sum_t \nabla_\theta \log \pi_\theta(\mathbf{a}_t^{(i)}|\mathbf{s}_t^{(i)}) \underbrace{\left(\sum_{t'} r(\mathbf{s}_{t'}^{(i)}, \mathbf{a}_{t'}^{(i)}) \right)}_{\hat{Q}_t^{(i)}}$
 3. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

$$\nabla J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla \log \pi_\theta(\mathbf{a}_t^{(i)}|\mathbf{s}_t^{(i)}) \underbrace{\hat{Q}_t^{(i)}}_{\text{"reward to go"}}$$

“reward to go”

$$\hat{Q}_t^{(i)} = \sum_{t'=t}^T r(\mathbf{s}_{t'}^{(i)}, \mathbf{a}_{t'}^{(i)})$$



Improving the policy gradient

$$\nabla J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla \log \pi_{\theta}(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \hat{Q}_t^{(i)}$$


$$\hat{Q}_t^{(i)} = \sum_{t'=t}^T r(\mathbf{s}_{t'}^{(i)}, \mathbf{a}_{t'}^{(i)})$$

$\hat{Q}_t^{(i)}$: estimate of expected reward if we take action $\mathbf{a}_t^{(i)}$ in state $\mathbf{s}_t^{(i)}$

can we get a better estimate?

$Q(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^T \mathbb{E}_{\pi_{\theta}} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t, \mathbf{a}_t]$: true *expected* reward-to-go

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) Q(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)})$$



What about the baseline?


$Q(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^T \mathbb{E}_{\pi_\theta} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t, \mathbf{a}_t]$: true *expected* reward-to-go

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_\theta \log \pi_\theta(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \left(Q(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) - V(\mathbf{s}_t^{(i)}) \right)$$

$$A(\mathbf{s}_t, \mathbf{a}_t) = Q(\mathbf{s}_t, \mathbf{a}_t) - V(\mathbf{s}_t)$$

b = average reward

$$V(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi_\theta(\mathbf{a}_t | \mathbf{s}_t)} [Q(\mathbf{s}_t, \mathbf{a}_t)]$$



State & state-action value functions

$Q^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^H E_{\pi_\theta} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t, \mathbf{a}_t]$: total reward from taking $\mathbf{a}_t$ in $\mathbf{s}_t$

$V^\pi(\mathbf{s}_t) = E_{\mathbf{a}_t \sim \pi_\theta(\mathbf{a}_t | \mathbf{s}_t)} [Q^\pi(\mathbf{s}_t, \mathbf{a}_t)]$: total reward from $\mathbf{s}_t$


$A^\pi(\mathbf{s}_t, \mathbf{a}_t) = Q^\pi(\mathbf{s}_t, \mathbf{a}_t) - V^\pi(\mathbf{s}_t)$: how much better $\mathbf{a}_t$ is

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_\theta \log \pi_\theta(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) A^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)})$$

the better this estimate, the lower the variance

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_\theta \log \pi_\theta(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \left(\sum_{t'=1}^H r(\mathbf{s}_{t'}^{(i)}, \mathbf{a}_{t'}^{(i)}) - b_t \right)$$

unbiased, but high variance single-sample estimate



Part 2:

Policy evaluation

Value function fitting

$$Q^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^H E_{\pi_\theta} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t, \mathbf{a}_t]$$

$$V^\pi(\mathbf{s}_t) = E_{\mathbf{a}_t \sim \pi_\theta(\mathbf{a}_t | \mathbf{s}_t)} [Q^\pi(\mathbf{s}_t, \mathbf{a}_t)]$$

$$A^\pi(\mathbf{s}_t, \mathbf{a}_t) = Q^\pi(\mathbf{s}_t, \mathbf{a}_t) - V^\pi(\mathbf{s}_t)$$

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_\theta \log \pi_\theta(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) A^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)})$$



fit *what* to *what*?

Q^π, V^π, A^π ?

$$Q^\pi(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + E_{\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)} [V^\pi(\mathbf{s}_{t+1})]$$

$$A^\pi(\mathbf{s}_t, \mathbf{a}_t) \approx r(\mathbf{s}_t, \mathbf{a}_t) + V^\pi(\mathbf{s}_{t+1}) - V^\pi(\mathbf{s}_t)$$

let's just fit $V^\pi(\mathbf{s})$!



Policy evaluation

$$V^\pi(\mathbf{s}_t) = \sum_{t'=t}^H E_{\pi_\theta} [r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t]$$

$$J(\theta) = E_{\mathbf{s}_1 \sim p(\mathbf{s}_1)} [V^\pi(\mathbf{s}_1)]$$



how can we perform policy evaluation?

Monte Carlo policy evaluation (this is what policy gradient does)

$$V^\pi(\mathbf{s}_t) \approx \sum_{t'=t}^H r(\mathbf{s}_{t'}, \mathbf{a}_{t'})$$

$$V^\pi(\mathbf{s}_t) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t'=t}^H r(\mathbf{s}_{t'}, \mathbf{a}_{t'})$$

(requires us to reset the simulator)



Monte Carlo evaluation with function approximation



$$V^\pi(\mathbf{s}_t) \approx \sum_{t'=t}^H r(\mathbf{s}_{t'}, \mathbf{a}_{t'})$$

not as good as this: $V^\pi(\mathbf{s}_t) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t'=t}^H r(\mathbf{s}_{t'}, \mathbf{a}_{t'})$

but still pretty good!

training data: $\left\{ \underbrace{\left(\mathbf{s}_t^{(i)}, \sum_{t'=t}^H r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \right)}_{y_t^{(i)}} \right\}$

$$\text{supervised regression: } \mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \sum_{t=1}^H \left\| \hat{V}_\phi^\pi(\mathbf{s}_t^{(i)}) - y_t^{(i)} \right\|^2$$



Can we do better?

ideal target: $y_t^{(i)} = \sum_{t'=t}^H E_{\pi_\theta} \left[r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) | \mathbf{s}_t^{(i)} \right] \approx r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + V^\pi(\mathbf{s}_{t+1}^{(i)}) \approx r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \underbrace{\hat{V}_\phi^\pi(\mathbf{s}_{t+1}^{(i)})}$

Monte Carlo target: $y_t^{(i)} = \sum_{t'=t}^H r(\mathbf{s}_{t'}, \mathbf{a}_{t'}^{(i)})$

directly use previous fitted value function!

training data: $\left\{ \underbrace{\left(\mathbf{s}_t^{(i)}, r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \hat{V}_\phi^\pi(\mathbf{s}_{t+1}^{(i)}) \right)}_{y_t^{(i)}} \right\}$

supervised regression: $\mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \sum_{t=1}^H \left\| \hat{V}_\phi^\pi(\mathbf{s}_t^{(i)}) - y_t^{(i)} \right\|^2$

sometimes referred to as a “bootstrapped” estimate

Discount factors

$$y_t^{(i)} = r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \hat{V}_\phi^\pi(\mathbf{s}_{t+1}^{(i)})$$

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \sum_{t=1}^H \left\| \hat{V}_\phi^\pi(\mathbf{s}_t^{(i)}) - y_t^{(i)} \right\|^2$$

what if H (episode length) is ∞ ?


$\hat{V}_\phi^\pi$ can get infinitely large in many cases

simple trick: better to get rewards sooner than later

$$y_t^{(i)} = r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \gamma \hat{V}_\phi^\pi(\mathbf{s}_{t+1}^{(i)})$$


↑
discount factor $\gamma \in [0, 1]$ (0.99 works well)

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[\sum_{t=1}^H \gamma^t r(\mathbf{s}_t, \mathbf{a}_t) \right]$$




episodic tasks

Iteration 2000



continuous/cyclical tasks

γ changes the MDP:





Time-varying vs. time-invariant

training data: $\left\{ \left(\mathbf{s}_t^{(i)}, r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \gamma \hat{V}_\phi^\pi(\mathbf{s}_{t+1}^{(i)}) \right) \right\}$

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \sum_{t=1}^H \left\| \hat{V}_\phi^\pi(\mathbf{s}_t^{(i)}) - y_t^{(i)} \right\|^2$$

we often don't actually care at which time step the data was collected

why?



training data: $\left\{ \left(\mathbf{s}_i, r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) \right) \right\}$

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$$

these are just notational conveniences,
we didn't actually change **anything**

Policy evaluation examples

TD-Gammon, Gerald Tesauro 1992




Figure 2. An illustration of the normal opening position in backgammon. TD-Gammon has sparked a near-universal conversion in the way experts play certain opening rolls. For example, with an opening roll of 4-1, most players have now switched from the traditional move of 13-9, 6-5, to TD-Gammon's preference, 13-9, 24-23. TD-Gammon's analysis is given in Table 2.




Figure 1. An illustration of the multilayer perception architecture used in TD-Gammon's neural network. This architecture is also used in the popular backpropagation learning procedure. Figure reproduced from [9].

AlphaGo, Silver et al. 2016



reward: game outcome

value function $\hat{V}_{\phi}^{\pi}(\mathbf{s}_t)$:

expected outcome given board state

reward: game outcome

value function $\hat{V}_{\phi}^{\pi}(\mathbf{s}_t)$:

expected outcome given board state




Part 3:

A basic actor-critic

A basic actor-critic algorithm


batch actor-critic algorithm:

1. sample $\{\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i\}$ from $\pi_\theta(\mathbf{a}|\mathbf{s})$ (run the policy)
2. evaluate $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$
3. refit $\hat{V}_\phi^\pi(\mathbf{s})$ to targets $\{y_i\}$
4. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
5. $\nabla_\theta J(\theta) \approx \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
6. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$




An online actor-critic algorithm

batch actor-critic algorithm:

- 
1. sample $\{\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i\}$ from $\pi_\theta(\mathbf{a}|\mathbf{s})$ (run the policy)
 2. evaluate $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$
 3. refit $\hat{V}_\phi^\pi(\mathbf{s})$ to targets $\{y_i\}$
 4. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
 5. $\nabla_\theta J(\theta) \approx \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
 6. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

what if we only have **one** transition?



$$\mathcal{L}_i(\phi) = \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$$

$$\phi \leftarrow \phi - \alpha \nabla_\phi \mathcal{L}_i(\phi)$$


$$\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$$

$$\nabla_\theta J(\theta) \approx \nabla_\theta \log \pi_\theta(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$$

$$\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$$


An online actor-critic algorithm

online actor-critic algorithm:

- 
1. get $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)$ by taking one step with $\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s})$
 2. evaluate $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$
 3. update ϕ using $\nabla_\phi \mathcal{L}_i(\phi)$
 4. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
 5. $\nabla_\theta J(\theta) \approx \nabla_\theta \log \pi_\theta(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
 6. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Example algorithm: A3C

Getting this to work in practice typically requires multiple parallel workers



$$\mathcal{L}_i(\phi) = \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$$

$$\phi \leftarrow \phi - \alpha \nabla_\phi \mathcal{L}_i(\phi)$$

$$\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$$

$$\nabla_\theta J(\theta) \approx \nabla_\theta \log \pi_\theta(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$$

$$\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$$

Intermission


Part 4:


On-policy actor-critic

(a.k.a. policy gradient that actually works)

A basic actor-critic algorithm

batch actor-critic algorithm:

- 
1. sample $\{\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i\}$ from $\pi_\theta(\mathbf{a}|\mathbf{s})$ (run the policy)
 2. evaluate $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$
 3. refit $\hat{V}_\phi^\pi(\mathbf{s})$ to targets $\{y_i\}$
 4. evaluate $\hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i) = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i) - \hat{V}_\phi^\pi(\mathbf{s}_i)$
 5. $\nabla_\theta J(\theta) \approx \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i|\mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i)$
 6. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$



What do we need to do to make this work in practice?

even better ways to estimate $\hat{A}^\pi$

e.g., GAE (today)

better ways to estimate $\nabla_\theta J(\theta)$

e.g., PPO (more on this later)

Critics and baselines

Actor-critic:
$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \left(r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \gamma \hat{V}_{\phi}^{\pi}(\mathbf{s}_{t+1}^{(i)}) - \hat{V}_{\phi}^{\pi}(\mathbf{s}_t^{(i)}) \right)$$

+ lower variance (due to critic)
- not unbiased (if the critic is not perfect)

Policy gradient:
$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \left(\left(\sum_{t'=t}^H \gamma^{t'-t} r(\mathbf{s}_{t'}^{(i)}, \mathbf{a}_{t'}^{(i)}) \right) - b_t \right)$$

+ no bias
- higher variance (because single-sample estimate)

can we use $\hat{V}_{\phi}^{\pi}$ and still keep the estimator unbiased?

Exercise: prove that this is still unbiased

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^H \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \left(\left(\sum_{t'=t}^H \gamma^{t'-t} r(\mathbf{s}_{t'}^{(i)}, \mathbf{a}_{t'}^{(i)}) \right) - \hat{V}_{\phi}^{\pi}(\mathbf{s}_t^{(i)}) \right)$$

+ no bias

+ lower variance (baseline is closer to rewards)

Eligibility traces & n-step returns

$$\hat{A}_C^\pi(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \hat{V}_\phi^\pi(\mathbf{s}_{t+1}) - \hat{V}_\phi^\pi(\mathbf{s}_t)$$

+ lower variance


- higher bias if value is wrong (it always is)

$$\hat{A}_{MC}^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^{\infty} \gamma^{t'-t} r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) - \hat{V}_\phi^\pi(\mathbf{s}_t)$$

+ no bias

- higher variance (because single-sample estimate)

Can we combine these two, to control bias/variance tradeoff?




$$\hat{A}_n^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) - \hat{V}_\phi^\pi(\mathbf{s}_t) + \gamma^n \hat{V}_\phi^\pi(\mathbf{s}_{t+n})$$

choosing $n > 1$ often works better!


bigger variance

cut here before variance gets too big!

smaller variance



Generalized advantage estimation



Do we have to choose just one n ?

Cut everywhere all at once!

$$\hat{A}_n^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) - \hat{V}_\phi^\pi(\mathbf{s}_t) + \gamma^n \hat{V}_\phi^\pi(\mathbf{s}_{t+n})$$

$$\hat{A}_{\text{GAE}}^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{n=1}^{\infty} w_n \hat{A}_n^\pi(\mathbf{s}_t, \mathbf{a}_t)$$

weighted combination of n -step returns

How to weight?

Mostly prefer cutting earlier (less variance)

$w_n \propto \lambda^{n-1}$ exponential falloff


$$\hat{A}_{\text{GAE}}^\pi(\mathbf{s}_t, \mathbf{a}_t) = r(\mathbf{s}_t, \mathbf{a}_t) + \gamma((1 - \lambda)\hat{V}_\phi^\pi(\mathbf{s}_{t+1}) + \lambda(r(\mathbf{s}_{t+1}, \mathbf{a}_{t+1}) + \gamma((1 - \lambda)\hat{V}_\phi^\pi(\mathbf{s}_{t+2}) + \lambda r(\mathbf{s}_{t+2}, \mathbf{a}_{t+2}) + \dots))$$

$$\hat{A}_{\text{GAE}}^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^{\infty} (\gamma\lambda)^{t'-t} \delta_{t'} \quad \delta_{t'} = r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) + \gamma\hat{V}_\phi^\pi(\mathbf{s}_{t'+1}) - \hat{V}_\phi^\pi(\mathbf{s}_{t'})$$

← similar effect as discount!

Policy gradient with GAE

policy gradient with GAE:

- 
1. sample $\{\tau^{(i)}\}$ from π_θ (run the policy)
 2. evaluate $y_t^{(i)} = r(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) + \gamma \hat{V}_\phi^\pi(\mathbf{s}_{t+1}^{(i)})$
 3. fit $\hat{V}_\phi^\pi(\mathbf{s})$ to targets $\{y_t^{(i)}\}$
 4. evaluate $\hat{A}_{\text{GAE}}^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) = \sum_{t'=t}^{\infty} (\gamma\lambda)^{t'-t} \delta_{t'}^{(i)}$
 5. $\nabla_\theta J(\theta) \approx \sum_i \sum_t \nabla_\theta \log \pi_\theta(\mathbf{a}_t^{(i)} | \mathbf{s}_t^{(i)}) \hat{A}^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)})$
 6. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

can use GAE estimator here too!




Other good options:

use n -step return estimators

use Monte Carlo estimator with $\hat{V}^\pi(\mathbf{s}_t)$ baseline

in practice, we typically
center the advantages



$$\mu = \frac{1}{HN} \sum_{i,t} \hat{A}^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)})$$

$$\sigma = \sqrt{\frac{1}{HN} \sum_{i,t} \left(\hat{A}^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) - \mu \right)^2}$$

$$\bar{A}^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) = \frac{\hat{A}^\pi(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}) - \mu}{\sigma}$$

This is the first part of implementing PPO

The other part is a slightly different gradient estimator


Part 5:

Off-policy actor-critic

(a.k.a. actor-critic)

Actor-critic with a replay buffer

online actor-critic algorithm:



not the action π_θ would have taken!

This algorithm is broken!

Can you spot the problems?

Hint: $(\mathbf{s}_i, \mathbf{a}_i)$ no longer comes from latest $\pi_\theta(\mathbf{a}_i|\mathbf{s}_i)$

Key idea: estimate Q-function

$$y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \hat{V}_\phi^\pi(\mathbf{s}'_i)$$

where $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)$ came from an older policy $\bar{\pi}$

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \left\| \hat{V}_\phi^\pi(\mathbf{s}_i) - y_i \right\|^2$$

why is this wrong?

↑

$$\text{want } V^\pi(\mathbf{s}_i) = \mathbb{E}_{\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s}_i)} [r(\mathbf{s}_i, \mathbf{a}) + \gamma V^\pi(\mathbf{s}')]]$$

$$\text{got } V^\pi(\mathbf{s}_i) = \mathbb{E}_{\mathbf{a} \sim \bar{\pi}(\mathbf{a}|\mathbf{s}_i)} [r(\mathbf{s}_i, \mathbf{a}) + \gamma V^\pi(\mathbf{s}')]]$$

$$\text{because } \mathbf{a}_i \sim \bar{\pi}(\mathbf{a}|\mathbf{s}_i)$$

What if we estimate Q-function instead?

$$y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \mathbb{E}_{\mathbf{a}' \sim \pi_\theta(\mathbf{a}'|\mathbf{s}'_i)} [\hat{Q}_\phi^\pi(\mathbf{s}'_i, \mathbf{a}')]]$$

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_{i=1}^N \left\| \hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a}_i) - y_i \right\|^2$$

to estimate:


$$\mathbf{a}'_i \sim \pi_\theta(\mathbf{a}'|\mathbf{s}'_i)$$

Note that we use the **latest** policy!

$$\mathbb{E}_{\mathbf{a}' \sim \pi_\theta(\mathbf{a}'|\mathbf{s}'_i)} [\hat{Q}_\phi^\pi(\mathbf{s}'_i, \mathbf{a}')]] \approx \hat{Q}_\phi^\pi(\mathbf{s}'_i, \mathbf{a}'_i)$$

Actor-critic with a replay buffer

online actor-critic algorithm:



use the same trick, but this time for $\mathbf{a}_i$ rather than $\mathbf{a}'_i$!

sample $\mathbf{a}_i^\pi \sim \pi_\theta(\mathbf{a}|\mathbf{s}_i)$

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i^\pi | \mathbf{s}_i) \hat{A}^\pi(\mathbf{s}_i, \mathbf{a}_i^\pi)$$


↑
not from replay buffer $\mathcal{R}$!

higher variance, but convenient
why is higher variance OK here?

in practice: $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\mathbf{a}_i^\pi | \mathbf{s}_i) \hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a}_i^\pi)$

Actor-critic with a replay buffer

off-policy actor-critic algorithm:



Is there any remaining problem?

$\mathbf{s}_i$ didn't come from $p_\theta(\mathbf{s})$


nothing we can do here, just accept it

intuition: we want optimal policy on $p_\theta(\mathbf{s})$

but we get optimal policy on a *broader* distribution

An implementation trick

off-policy actor-critic algorithm:

- 
1. get $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)$ by taking one step with $\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s})$
 2. evaluate $y_i = r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \mathbb{E}_{\mathbf{a}' \sim \pi_\theta(\mathbf{a}'|\mathbf{s}'_i)}[\hat{Q}_\phi^\pi(\mathbf{s}'_i, \mathbf{a}')]]$
 3. update ϕ using $\nabla_\phi \sum_{i=1}^B \|\hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a}_i) - y_i\|^2$
 4. $\nabla_\theta J(\theta) \approx \sum_i \mathbb{E}_{\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s}_i)}[\nabla_\theta \log \pi_\theta(\mathbf{a}|\mathbf{s}_i) \hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a})]$
 5. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

which gradient is this estimating?

$$\nabla_\theta \mathbb{E}_{\mathbf{a} \sim \pi_\theta(\mathbf{a}|\mathbf{s}_i)}[\hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a})]$$

this is a neural network!

we can compute $\frac{d}{d\mathbf{a}} \hat{Q}_\phi^\pi(\mathbf{s}_i, \mathbf{a})$

can we use this?

Reparameterization trick

say $\pi_{\theta}(\mathbf{a}|\mathbf{s}) = \mathcal{N}(\mu_{\theta}(\mathbf{s}), \sigma_{\theta}(\mathbf{s}))$

how do we sample?

$$\epsilon \sim \mathcal{N}(0, 1)$$

$$\mathbf{a} = \mu_{\theta}(\mathbf{s}) + \sigma_{\theta}(\mathbf{s})\epsilon$$

$$\nabla_{\theta} \mathbb{E}_{\mathbf{a} \sim \pi_{\theta}(\mathbf{a}|\mathbf{s})} [Q(\mathbf{s}, \mathbf{a})] \approx \nabla_{\theta} Q(\mathbf{s}, \mu_{\theta}(\mathbf{s}) + \sigma_{\theta}(\mathbf{s})\epsilon)$$

“reparameterized gradient”


$$\mathbb{E}_{\mathbf{a} \sim \pi_{\theta}(\mathbf{a}|\mathbf{s})} [Q(\mathbf{s}, \mathbf{a})] = \mathbb{E}_{\epsilon \sim \mathcal{N}(0, 1)} [Q(\mathbf{s}, \mu_{\theta}(\mathbf{s}) + \sigma_{\theta}(\mathbf{s})\epsilon)]$$


$$\approx Q(\mathbf{s}, \mu_{\theta}(\mathbf{s}) + \sigma_{\theta}(\mathbf{s})\epsilon) \quad \epsilon \sim \mathcal{N}(0, 1)$$



this is easy to differentiate with any
autograd library (e.g., PyTorch)

Off-policy actor-critic with reparameterization

off-policy reparameterized actor-critic algorithm:



Example practical algorithm:

Soft actor-critic

- uses a few more tricks for Q-function estimation (we'll see this in Q-learning lectures)
- uses entropy regularization (we'll learn about this later)

We'll also learn about algorithms that do this with deterministic policies later!