



10-x23 Generative AI

Machine Learning Department
School of Computer Science
Carnegie Mellon University

Long Context in LLMs

Matt Gormley
Lecture 19

Reminders

- **Programming Test: HW3 / HW4**
 - Date: Friday, Mar 27
 - Time: 40 minutes
 - Check Piazza for seat assignment
- **Exam**
 - Date: evening exam, Monday, Mar 30 at 7:00pm – 8:20pm
 - Time: 80 minutes
 - Covered Material: Lectures 1 – 15 (same as Quiz 1 – Quiz 4)
 - You may bring one sheet of notes (front and back)
 - Format of questions: Unlike the Quiz questions, which were all multiple choice, Exam questions will include open-ended questions as well
 - Check Piazza for seat assignment

LONG-CONTEXT LLMS

What can I help with?

Write a python script that trains a decoder-only Transformer model for video and text understanding.

1. The model should start with a pretrained large language model (LLM) as the backbone model.
2. Use a pretrained ViT model. Convert each video to 224x224 resolution at just 10 frames per second. Each frame should then be encoded using a Vision Transformer (ViT) to 66 tokens (i.e. one token per 14x14 patch). These image tokens should be delineated by special tokens <frame_n> </frame_n> where n indicates the frame number.
3. The model should be fine-tuned on a video question-answering dataset.



Search



Deep research



```

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
import torchvision.io as io
from transformers import AutoTokenizer, AutoModelForCausalLM
import timm
import os

# Special tokens
def frame_tokens(n):
    return f"<frame_{n}>", f"</frame_{n}>"

# 1. ViT Encoder
class ViTFrameEncoder(nn.Module):
    def __init__(self, model_name='vit_base_patch14_224', device='cuda'):
        super().__init__()
        self.vit = timm.create_model(
            model_name, pretrained=True, num_classes=0).to(device)
        self.device = device
        self.transform = transforms.Compose([
            transforms.Resize((224, 224)),
            transforms.ToTensor(),
            transforms.Normalize([0.5]*3, [0.5]*3)
        ])

    def forward(self, video_path, fps=10, max_frames=10):
        video, _, _ = io.read_video(video_path, pts_unit='sec')
        frame_indices = torch.linspace(0, video.shape[0]-1, steps=max_frames).long()
        frames = video[frame_indices]
        frames = torch.stack(
            [self.transform(frame.permute(2, 0, 1).float() / 255.) for frame in frames])
        B, C, H, W = frames.shape
        frames = frames.to(self.device)
        with torch.no_grad():
            # (B, tokens, D)
            features = self.vit.forward_features(frames)
        return features # shape: (frames, 66, D)

```

```

# 2. Dataset
class VideoQADataset(Dataset):
    def __init__(self, data, tokenizer, frame_encoder, max_length=512):
        self.data = data
        self.tokenizer = tokenizer
        self.frame_encoder = frame_encoder
        self.max_length = max_length

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        item = self.data[idx]
        video_path = item['video']
        question = item['question']
        answer = item['answer']

        image_tokens = []
        features = self.frame_encoder(video_path)
        for i, frame_tokens_embed in enumerate(features):
            start_tok, end_tok = frame_tokens(i)
            image_tokens.append(start_tok)
            # use dummy placeholders
            image_tokens.extend(["<img_tok>"] * frame_tokens_embed.shape[0])
            image_tokens.append(end_tok)

        text = ' '.join(image_tokens) + ' ' + question
        inputs = self.tokenizer(text, return_tensors='pt',
                                truncation=True, max_length=self.max_length)
        labels = self.tokenizer(answer, return_tensors='pt',
                                truncation=True, max_length=self.max_length) ['input_ids']

        return {
            'input_ids': inputs['input_ids'].squeeze(0),
            'attention_mask': inputs['attention_mask'].squeeze(0),
            'labels': labels.squeeze(0)
        }

```

```

# 3. Decoder Model
class VideoTextDecoder(nn.Module):
    def __init__(self, llm_name='gpt2', image_token_dim=768, device='cuda'):
        super().__init__()
        self.tokenizer = AutoTokenizer.from_pretrained(llm_name)
        self.tokenizer.add_special_tokens({'additional_special_tokens': [
            f"<frame_{i}>" for i in range(10)] +
        [f"</frame_{i}>" for i in range(10)] + ["<img_tok>"]})
        self.model = AutoModelForCausalLM.from_pretrained(llm_name).to(device)
        self.model.resize_token_embeddings(len(self.tokenizer))
        self.device = device
        self.image_token_dim = image_token_dim

    def forward(self, input_ids, attention_mask, labels):
        outputs = self.model(input_ids=input_ids,
                              attention_mask=attention_mask, labels=labels)

        return outputs

# 4. Training Loop
def train(model, dataset, tokenizer, epochs=3, batch_size=2, lr=5e-5):
    dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
    optimizer = torch.optim.AdamW(model.parameters(), lr=lr)

    for epoch in range(epochs):
        model.train()
        for batch in dataloader:
            input_ids = batch['input_ids'].to(model.device)
            attention_mask = batch['attention_mask'].to(model.device)
            labels = batch['labels'].to(model.device)

            outputs = model(input_ids, attention_mask, labels)
            loss = outputs.loss
            loss.backward()
            optimizer.step()
            optimizer.zero_grad()
            print(f"Epoch {epoch}, Loss: {loss.item():.4f}")

```

```
# Example Usage
if __name__ == "__main__":
    device = 'cuda' if torch.cuda.is_available() else 'cpu'
    vit_encoder = ViTFrameEncoder(device=device)
    decoder = VideoTextDecoder(device=device)

    # Replace this with actual dataset
    mock_data = [
        {'video': 'sample_video.mp4', 'question': 'What is the person doing?',
         'answer': 'Running on a beach.'}
    ]
    dataset = VideoQADataset(mock_data, decoder.tokenizer, vit_encoder)

    train(decoder, dataset, decoder.tokenizer)
```


Using this model, how many tokens will be produced by encoding 3 seconds of video. Show your work.

Step-by-step Breakdown:

1. Frame Sampling Rate:

- 10 frames per second
- For 3 seconds →
 $10 \text{ fps} \times 3 \text{ seconds} = 30 \text{ frames}$

2. ViT Tokenization:

- Each frame is resized to **224x224**
- Patch size is **14x14** →
 $\frac{224}{14} = 16$ patches per side
 $16 \times 16 = 256$ patches → **256 tokens per frame**

3. Special Tokens per Frame:

- Each frame is wrapped in:
"`<frame_n>`" and "`</frame_n>`" → **2 special tokens**

4. Total Tokens per Frame:

- $256 \text{ (ViT tokens)} + 2 \text{ (special tokens)} = 258 \text{ tokens}$

5. Total Tokens for 30 Frames:

- $30 \text{ frames} \times 258 \text{ tokens/frame} = \boxed{7740 \text{ tokens}}$

How about 30 minutes of video?

Final Answer:

30 minutes of video → 4,644,000 tokens

That's **way beyond** the context window of most LLMs (usually 2k to 32k tokens), so you'd need to:

- Chunk the video (e.g., per scene or question)
- Use hierarchical attention
- Compress image tokens (e.g., pooling, vector quantization)
- Use retrieval or memory-based attention mechanisms

Want help with any of those strategies?

Context Length of Transformer LMs

Comparison of some recent **large language models** (LLMs) on their **context size**, i.e. how many tokens they can accept

Model	Creators	Year	Model Size	Context Size
GPT-2	OpenAI	2019	1.5 billion	1024
GPT-3	OpenAI	2020	175 billion	2048
PaLM	Google	2022	540 billion	2048
LLaMA	Meta	2023	65 billion	2048
LLaMA-2	Meta	2023	70 billion	4096
Claude-2	Anthropic	2023	? (130 billion)	100k
Claude-2.1	Anthropic	2023	? (130 billion)	200k
GPT-4	OpenAI	2023	? (1.76 trillion)	8192
Mistral	Mistral AI	2023	7 billion	8192 (32k)
Mixtral	Mistral AI	2023	47 billion	8192 (128k)
Gemini (Ultra)	Google	2023	? (1.5 trillion)	32k
LWModel	academia!	2023	7 billion	1 million
Gemini-1.5	Google	2024	? (1.5 trillion)	1 million
GPT-4o	OpenAI	2024	?	128k
LLaMA-3	Meta	2024	405 billion	8192
LLaMA-3.1	Meta	2024	405 billion	128k
Claude-3.5	Anthropic	2024	?	200k


Context Length of Transformer LMs

Comparison of some recent **large language models** (LLMs) on their **context size**, i.e. how many tokens they can accept

Model	Creators	Year	Model Size	Context Size
Qwen2.5	Alibaba / Qwen	2024	72B	128k
GPT-4.1	OpenAI	2025	?	1M
Gemini 2.5 Pro	Google	2025	?	1M
Llama 4 Scout	Meta	2025	17B active / 109B total	10M
Llama 4 Maverick	Meta	2025	17B active / 400B total	10M
Kimi K2	Moonshot AI	2025	1T total / 32B active	256k
Qwen3	Alibaba / Qwen	2025	235B total / 22B active	32k
GPT-5	OpenAI	2025	?	400k
GPT-5.4	OpenAI	2026	?	1M
Gemini 3.1 Pro	Google	2026	?	1M
Claude Sonnet 4.6 / Opus 4.6	Anthropic	2026	?	1M
Kimi K2.5	Moonshot AI	2026	1T total / 32B active	256k
Qwen3.5-Plus	Alibaba / Qwen	2026	?	1M
Grok 4.20	xAI	2026	?	2M


Pressure Testing Claude-2.1 200K via "Needle In A HayStack"

Asking Claude 2.1 To Do Fact Retrieval Across Context Lengths & Document Depth



Asking Claude 2.1 To Do Fact Retrieval Across Context Lengths & Document Depth


Asking Claude 2.1 To Do Fact Retrieval Across Context Lengths & Document Depth



A fact was placed within a document. Claude 2.1 (200K) was then asked to retrieve it. The output was evaluated (with GPT-4) for accuracy. This test was run at 35 different document depths (top > bottom) and 35 different context lengths (1K > 200K tokens). Document Depths followed a sigmoid distribution

Extending Short-Context Models

- There are two key ingredients for extending a short context model
 1. extensible positional embeddings
 2. careful selection of long data
- General recipe
 - Pre-train a short context model (block size = 4k) on 1 trillion tokens of text
 - Adjust the hyperparameters of the positional embeddings
 - Continue pre-training but increase the block size to support long-contexts (block size = 80k) on only 5 billion tokens of text



Fine-Tuning vs. In-Context Learning

- Why would we ever bother with fine-tuning if it's so inefficient?
- Because, even for very large LMs, fine-tuning often beats in-context learning
- In a fair comparison of fine-tuning (FT) and in-context learning (ICL), we find that FT outperforms ICL for most model sizes on RTE and MNLI

		FT						
		125M	350M	1.3B	2.7B	6.7B	13B	30B
ICL	125M	-0.00	0.01	0.02	0.03	0.12	0.14	0.09
	350M	-0.00	0.01	0.02	0.03	0.12	0.14	0.09
	1.3B	-0.00	0.01	0.02	0.03	0.12	0.14	0.09
	2.7B	-0.00	0.01	0.02	0.03	0.12	0.14	0.09
	6.7B	-0.00	0.01	0.02	0.03	0.12	0.14	0.09
	13B	-0.04	-0.02	-0.01	-0.00	0.09	0.11	0.05
	30B	-0.11	-0.09	-0.08	-0.08	0.02	0.03	-0.02

(a) RTE

		FT						
		125M	350M	1.3B	2.7B	6.7B	13B	30B
ICL	125M	-0.00	0.00	0.02	0.01	0.10	0.11	0.07
	350M	-0.00	0.00	0.02	0.01	0.10	0.11	0.07
	1.3B	-0.01	-0.00	0.01	0.01	0.10	0.11	0.07
	2.7B	-0.01	-0.00	0.01	0.01	0.09	0.10	0.07
	6.7B	-0.01	-0.01	0.01	0.00	0.09	0.10	0.06
	13B	-0.03	-0.03	-0.02	-0.02	0.07	0.08	0.04
	30B	-0.07	-0.07	-0.05	-0.06	0.03	0.04	0.00

(b) MNLI


At least this was the general wisdom in 2023.

We might have a different story to tell now that it's 2025.
(See Lecture 19)


Table 1: Difference between average **out-of-domain performance** of ICL and FT on RTE (a) and MNLI (b) for different model sizes. We use 16 examples and 10 random seeds for both approaches. For ICL, we use the standard in-context learning prompt. For FT, we use pattern-based fine-tuning (PBFT) and select checkpoints according to in-domain performance. We perform a Welch's t-test and color cells according to whether: **ICL performs significantly better than FT** (red), **FT performs significantly better than ICL** (blue). For cells without color, there is no significant difference.

ICL with Long Sequences

- Modern wisdom reveals that ICL can sometimes outperform fine-tuning
- All we need is a long context model that can hold our in-context demonstrations
- Two multiclass classification datasets (Clinic-150 with 151 labels, Trecfine with 50 labels)
- Base model: LLaMa2-7B
- Approaches:
 - Finetuned: LoRA
 - Random ICL: randomly selects training examples for ICL
 - Retrieval ICL: selects training examples for ICL most similar to test example based on BM25



(a) Clinic-150



(b) Trecfine

APPROXIMATE ATTENTION

Approximate Attention

- Standard attention requires $O(N^2)$ memory and computation
- While the computation requirement may be acceptable, the memory requirement is usually not
- One solution is to instead approximate the attention computation
- Examples include:
 - Sparse Attention (2019), $O(N \sqrt{N})$
 - Sliding Window Attention (2020), $O(N)$
 - Dilated Attention (2023), $O(N)$

Sparse Attention

4.2. Factorized self-attention


A self-attention layer maps a matrix of input embeddings X to an output matrix and is parameterized by a connectivity pattern $S = \{S_1, \dots, S_n\}$, where S_i denotes the set of indices of the input vectors to which the i th output vector attends. The output vector is a weighted sum of transformations of the input vectors:

$$\text{Attend}(X, S) = \left(a(\mathbf{x}_i, S_i) \right)_{i \in \{1, \dots, n\}} \quad (2)$$

$$a(\mathbf{x}_i, S_i) = \text{softmax} \left(\frac{(W_q \mathbf{x}_i) K_{S_i}^T}{\sqrt{d}} \right) V_{S_i} \quad (3)$$

$$K_{S_i} = \left(W_k \mathbf{x}_j \right)_{j \in S_i} \quad V_{S_i} = \left(W_v \mathbf{x}_j \right)_{j \in S_i} \quad (4)$$

Here W_q , W_k , and W_v represent the weight matrices which transform a given $\mathbf{x}_i$ into a *query*, *key*, or *value*, and d is the inner dimension of the queries and keys. The output at each position is a sum of the values weighted by the scaled dot-product similarity of the keys and queries.



(a) Transformer

(c) Sparse Transformer (fixed)


Sliding Window Attention

Sliding Window Attention


- also called “local attention” and introduced for the Longformer model (2020)
- **The problem:** regular attention is computationally expensive and requires a lot of memory
- **The solution:** apply a causal mask that only looks at the include a window of $(\frac{1}{2}w+1)$ tokens, with the rightmost window element being the current token (i.e. on the diagonal)

$$\mathbf{X}' = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} + \mathbf{M} \right) \mathbf{V}$$


regular causal attention



sliding window attention (w=4)




sliding window attention (w=6)



Sliding Window Attention

Sliding Window Attention

- also called “local attention” and introduced for the Longformer model (2020)
- **The problem:** regular attention is computationally expensive and requires a lot of memory
- **The solution:** apply a causal mask that only looks at the include a window of $(\frac{1}{2}w+1)$ tokens, with the rightmost window element being the current token (i.e. on the diagonal)




3 ways you could implement


1. *naïve implementation:* just do the matrix multiplication, but this is still slow
2. *for-loop implementation:* asymptotically faster / less memory, but unusable in practice b/c for-loops in PyTorch are too slow
3. *sliding chunks implementation:* break into Q and K into chunks of size $w \times w$, with overlap of $\frac{1}{2}w$; then compute full attention within each chunk and mask out chunk (very fast/low memory in practice)

$$\mathbf{X}' = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} + \mathbf{M} \right) \mathbf{V}$$


Dilated Attention




Segment Length: 4
Dilated Rate: 1




Segment Length: 8
Dilated Rate: 2



Segment Length: 16
Dilated Rate: 4



- Dilated attention mixes together multiple dilation rates
- Each dilation rate decides the sparsity pattern of Q,K,V
- Computation is easily parallelizable
- Runtime is great, but it's now a different model altogether



EFFICIENT FULL ATTENTION

Scaling Up Computation with Context Length

- In TransformerLMs, the FLOPS do not scale up as quickly as you might expect with context size
- There are lots of computationally intensive components to the Transformer besides the $O(N^2)$ attention






Figure 5: The per dataset training FLOPs cost ratio relative to a 4k context size, considering different model dimensions. On the x-axis, you'll find the context length, where, for example, 32x(128k) denotes a context length of 128k, 32x the size of the same model's 4k context length.

Sequence Parallelism


- **Sequence parallelism** breaks apart a long sequence into chunks
- Each chunk is given to a separate device (GPU or TPU) and the computation of earlier devices must be sent to later devices (for decoder-only Transformers)
- **Problem:** this does not scale up very efficiently because the later queries still must attend to $O(N)$ other key/value tokens



(a) Pipeline parallelism




(b) Tensor parallelism



(c) Sequence parallelism (Ours)

Blockwise Parallel Transformer (BPT)

- FlashAttention provides a dramatic reduction in the Transformer's memory requirement, but is only does blockwise computation in the attention layers
- Blockwise Parallel Transformer (BPT) does blockwise computation of attention **and** the feed-forward neural network layer and residual connection



Blockwise Parallel Transformer (BPT)

Recall: softmax is shift invariant! So we can compute attention in blocks and rescale the prior outputs appropriately based on the sufficient statistics from the other blocks

This is achieved by keeping track of normalization statistics and combining them from all blocks to scale each block accordingly. For a specific query block Q_i , $1 \leq i \leq B_q$, the corresponding attention output can be computed by scaling each blockwise attention as follows:

$$\text{Attention}(Q_i, K, V) = \text{Scaling}(\{\exp(Q_i K_j^T) V_j\}_{j=1}^{B_{kv}}). \quad (3)$$

The scaling operation scales each blockwise attention based on the difference between the blockwise maximum and the global maximum:

$$\begin{aligned} \text{Attention}(Q_i, K_j, V_j) &= \exp(Q_i K_j^T - \max(Q_i K_j^T)) / \sum \exp(Q_i K_j^T - \max(Q_i K_j^T)) \\ \max_i &= \max(\max(Q_i K_1^T), \dots, \max(Q_i K_B^T)) \\ \text{Attention}(Q_i, K, V) &= [\exp(Q_i K_j^T - \max_i) \text{Attention}(Q_i, K_j, V_j)]_{j=1}^{B_{kv}}. \end{aligned}$$

This blockwise self-attention computation eliminates the need to materialize the full attention matrix of size $O(n^2)$, resulting in significant memory savings.

Blockwise Parallel Transformer (BPT)

Note: BPT computes the attention + FFN layers together for each query block Q_i by iterating over all key/value blocks K_j, V_j

We observe that the blockwise computation is not limited to self-attention but can also be applied to the feedforward network. For each query block, after iterating over the key and value blocks, the feedforward network can be computed along with a residual connection, completing the attention and feedforward network computation for that query block. This means that the model does not need to compute the feedforward network on the full sequence, but rather on intermediate blocks, resulting in memory savings. The computation for a query block is given by:

$$\text{Output}_i = \text{FFN}(\text{Attention}(Q_i, K, V) + Q_i) + \text{Attention}(Q_i, K, V) + Q_i.$$

Therefore, the output for each block consists of the feedforward network, self-attention, and residual connection computed in a blockwise manner.

Blockwise Parallel Transformer (BPT)

Algorithm 1 Reduce memory cost with BPT.

Required: Input sequence x . Number of query blocks B_q . Number of key and value blocks B_{kv} .

Initialize

Project input sequence x into query, key and value.

Split query sequence into B_q of query input blocks.

Split key and value sequences into B_{kv} of key-value input blocks.

for $outer = 1$ **to** B_q **do**

 Choose the $outer$ -th query.

for $inner = 1$ **to** B_{kv} **do**

 Choose the $inner$ -th key and $inner$ -th value block.

 Compute attention using query, key and value, and record normalization statistics.

end for

 Combine each blocks by scaling them to get attention output for the $outer$ -th input block.

 Compute feedforward on attention output and add residual connection.

end for

Blockwise Parallel Transformer (BPT)

Blockwise Parallel Transformer (BPT)

Table 2: Maximum context length during training with different methods. BPT enables training 2-4 times longer sequence length than FlashAttention / Memory Efficient Attention, and up to 32 times longer sequence length than vanilla attention.

1 A100	PartitionSpec	Vanilla Attention	MemoryEfficient	Blockwise Parallel
350M	(1,1,1)	16K (16384)	65K (65536)	131K (131072)
1B	(1,1,1)	16K (16384)	65K (65536)	131K (131072)
3B	(1,1,1)	8K (8192)	16K (16384)	65K (65536)
8 A100	PartitionSpec	Vanilla Attention	MemoryEfficient	Blockwise Parallel
3B	(1,1,8)	16K (16384)	65K (65536)	131K (131072)
7B	(1,1,8)	16K (16384)	65K (65536)	131K (131072)
13B	(1,1,8)	8K (8192)	33K (32768)	65K (65536)
30B	(1,1,8)	8K (8192)	16K (16384)	65K (65536)
64 TPUv4	PartitionSpec	Vanilla Attention	MemoryEfficient	Blockwise Parallel
13B	(1,1,64)	4K (4096)	16K (16384)	33K (32768)
30B	(1,1,64)	2K (2048)	4K (4096)	16K (16384)
70B	(1,1,64)	1k (1024)	2K (2048)	8K (8192)

Ring Attention


- Key idea:
 - use Blockwise Parallel Transformers (BPT) to process very long input sequences using many devices
 - communicate the key-value blocks between devices in exactly the same pattern that the key-value blocks are used in the attention computation
- Key outcome:
 - the sequence length that fits in memory (during training and inference) scales with the number of devices available

Ring Attention (aka. Pass the Parcel)




Ring Attention

Ring Attention



Ring Attention



Ring Attention

Algorithm 1 Reducing Transformers Memory Cost with Ring Attention.

Required: Input sequence x . Number of hosts N_h .

Initialize

Split input sequence into N_h blocks that each host has one input block.

Compute query, key, and value for its input block on each host.

for Each transformer layer **do**

for $count = 1$ **to** $N_h - 1$ **do**

for For each host concurrently. **do**

 Compute memory efficient attention incrementally using local query, key, value blocks.

 Send key and value blocks to next host and receive key and value blocks from previous host.

end for

end for

for For each host concurrently. **do**

 Compute memory efficient feedforward using local attention output.

end for

end for

Ring Attention

- Results:
 - 13B model can be increased to handle a 128k token context length on 8 A100s
 - 7B model can be increased to handle a 4 million token context length on 32 A100s

	Max context size supported ($\times 1e3$)				Ours vs SOTA
	Vanilla	Flash Attn	BPT	Ring Attention (Ours)	
8x A100 NVLink					
3B	4	32	64	512	8x
7B	2	16	32	256	8x
13B	2	4	16	128	8x
32x A100 InfiniBand					
7B	4	64	128	4096	32x
13B	4	32	64	2048	32x

Large World Model




Figure 1 LWM can answer questions over a 1 hour YouTube video. Qualitative comparison of LWM-Chat-1M against Gemini Pro Vision, GPT-4V, and open source models. Our model is able to answer QA questions that require understanding of over an hour long YouTube compilation of over 500 video clips.




Figure 2 LWM can retrieval facts across 1M context with high accuracy. Needle retrieval comparisons against Gemini Pro and GPT-4 for each respective max context length – 32K and 128K. Our model performs competitively while being able to extend to 8x longer context length. Note that in order to show fine-grained results, the x-axis is log-scale from 0-128K, and linear-scale from 128K-1M.